Assignment 2 - Matrix Multiplication
COMP30250 - Parallel and Cluster Computing

October 29, 2021

Aleryc SERRANIA - 21204068

1 Introduction

This report was created with a Jupyter notebook. You can find the original source of this notebook in the same folder (`report.ipynb`). Source code of the implementations of the algorithms can be found inside `src` folder. A Makefile is provided to compile the code.

The variants for this assignments are the following:

1. Manually written straightforward non-blocked ijk algorithm;
2. Blocked ijk algorithm using ATLAS calls to compute multiplication of blocks;
3. Blocked kij algorithm using ATLAS calls to compute multiplication of blocks;
4. Compare the fastest program with BLAS dgemm routine.

I built a python script `run_exp.py` to get results for different sizes of matrix and block and also to run multiple iterations of the same configuration to get an average of the execution time (25 iterations). The results are compiled inside `data.csv`.

2 Detailed performance results for each algorithm

Timing unit is second.

2.1 Straightforward non-blocked ijk

[3]:

	filename	matrix_size	timing
0	non_blocked_ijk.out	2	2.000000e-07
1	non_blocked_ijk.out	4	7.600000e-07
2	non_blocked_ijk.out	8	3.480000e-06
3	non_blocked_ijk.out	16	2.420000e-05
4	non_blocked_ijk.out	32	2.048400e-04
5	non_blocked_ijk.out	64	1.452720e-03
6	non_blocked_ijk.out	128	1.265856e-02
7	non_blocked_ijk.out	256	1.499471e-01
8	non_blocked_ijk.out	512	1.256371e+00
9	non_blocked_ijk.out	1024	2.541357e+01

2.2 Blocked ijk with ATLAS

[4]:

	filename	matrix_size	block_size	timing
20	blocked_ijk_atlas.out	2	2	0.000017
21	blocked_ijk_atlas.out	4	2	0.000020
22	blocked_ijk_atlas.out	4	4	0.000018
23	blocked_ijk_atlas.out	8	2	0.000026
24	blocked_ijk_atlas.out	8	4	0.000018
25	blocked_ijk_atlas.out	8	8	0.000017
26	blocked_ijk_atlas.out	16	2	0.000055
27	blocked_ijk_atlas.out	16	4	0.000026
28	blocked_ijk_atlas.out	16	8	0.000020
29	blocked_ijk_atlas.out	16	16	0.000021
30	blocked_ijk_atlas.out	32	2	0.000341
31	blocked_ijk_atlas.out	32	4	0.000081
32	blocked_ijk_atlas.out	32	8	0.000043
33	blocked_ijk_atlas.out	32	16	0.000034
34	blocked_ijk_atlas.out	32	32	0.000035
35	blocked_ijk_atlas.out	64	2	0.002313
36	blocked_ijk_atlas.out	64	4	0.000451
37	blocked_ijk_atlas.out	64	8	0.000183
38	blocked_ijk_atlas.out	64	16	0.000140
39	blocked_ijk_atlas.out	64	32	0.000140
40	blocked_ijk_atlas.out	64	64	0.000138
41	blocked_ijk_atlas.out	128	2	0.018559
42	blocked_ijk_atlas.out	128	4	0.003531
43	blocked_ijk_atlas.out	128	8	0.001394
44	blocked_ijk_atlas.out	128	16	0.001014
45	blocked_ijk_atlas.out	128	32	0.000939
46	blocked_ijk_atlas.out	128	64	0.000588
47	blocked_ijk_atlas.out	128	128	0.000648
48	blocked_ijk_atlas.out	256	2	0.163757
49	blocked_ijk_atlas.out	256	4	0.030790
50	blocked_ijk_atlas.out	256	8	0.011338
51	blocked_ijk_atlas.out	256	16	0.008187
52	blocked_ijk_atlas.out	256	32	0.007289
53	blocked_ijk_atlas.out	256	64	0.004172
54	blocked_ijk_atlas.out	256	128	0.004046
55	blocked_ijk_atlas.out	256	256	0.004101
56	blocked_ijk_atlas.out	512	2	1.307175
57	blocked_ijk_atlas.out	512	4	0.247788
58	blocked_ijk_atlas.out	512	8	0.093341
59	blocked_ijk_atlas.out	512	16	0.063817
60	blocked_ijk_atlas.out	512	32	0.058434
61	blocked_ijk_atlas.out	512	64	0.032584
62	blocked_ijk_atlas.out	512	128	0.030997
63	blocked_ijk_atlas.out	512	256	0.030732

64	blocked_ijk_atlas.out	512	512	0.029372
65	blocked_ijk_atlas.out	1024	2	13.861305
66	blocked_ijk_atlas.out	1024	4	3.207080
67	blocked_ijk_atlas.out	1024	8	0.949980
68	blocked_ijk_atlas.out	1024	16	0.530183
69	blocked_ijk_atlas.out	1024	32	0.478315
70	blocked_ijk_atlas.out	1024	64	0.268475
71	blocked_ijk_atlas.out	1024	128	0.251911
72	blocked_ijk_atlas.out	1024	256	0.245424
73	blocked_ijk_atlas.out	1024	512	0.228122
74	blocked_ijk_atlas.out	1024	1024	0.225221

2.3 Blocked kij with ATLAS

[5]:

	filename	matrix_size	block_size	timing
75	blocked_kij_atlas.out	2	2	0.000015
76	blocked_kij_atlas.out	4	2	0.000015
77	blocked_kij_atlas.out	4	4	0.000015
78	blocked_kij_atlas.out	8	2	0.000020
79	blocked_kij_atlas.out	8	4	0.000016
80	blocked_kij_atlas.out	8	8	0.000017
81	blocked_kij_atlas.out	16	2	0.000062
82	blocked_kij_atlas.out	16	4	0.000023
83	blocked_kij_atlas.out	16	8	0.000017
84	blocked_kij_atlas.out	16	16	0.000017
85	blocked_kij_atlas.out	32	2	0.000290
86	blocked_kij_atlas.out	32	4	0.000075
87	blocked_kij_atlas.out	32	8	0.000040
88	blocked_kij_atlas.out	32	16	0.000034
89	blocked_kij_atlas.out	32	32	0.000030
90	blocked_kij_atlas.out	64	2	0.002284
91	blocked_kij_atlas.out	64	4	0.000540
92	blocked_kij_atlas.out	64	8	0.000202
93	blocked_kij_atlas.out	64	16	0.000157
94	blocked_kij_atlas.out	64	32	0.000138
95	blocked_kij_atlas.out	64	64	0.000128
96	blocked_kij_atlas.out	128	2	0.017508
97	blocked_kij_atlas.out	128	4	0.003818
98	blocked_kij_atlas.out	128	8	0.001504
99	blocked_kij_atlas.out	128	16	0.001120
100	blocked_kij_atlas.out	128	32	0.000930
101	blocked_kij_atlas.out	128	64	0.000595
102	blocked_kij_atlas.out	128	128	0.000671
103	blocked_kij_atlas.out	256	2	0.138930
104	blocked_kij_atlas.out	256	4	0.028898
105	blocked_kij_atlas.out	256	8	0.011014
106	blocked_kij_atlas.out	256	16	0.007977

107	blocked_kij_atlas.out	256	32	0.007328
108	blocked_kij_atlas.out	256	64	0.004165
109	blocked_kij_atlas.out	256	128	0.004155
110	blocked_kij_atlas.out	256	256	0.004052
111	blocked_kij_atlas.out	512	2	1.107016
112	blocked_kij_atlas.out	512	4	0.228654
113	blocked_kij_atlas.out	512	8	0.087474
114	blocked_kij_atlas.out	512	16	0.063513
115	blocked_kij_atlas.out	512	32	0.057866
116	blocked_kij_atlas.out	512	64	0.032614
117	blocked_kij_atlas.out	512	128	0.030717
118	blocked_kij_atlas.out	512	256	0.030706
119	blocked_kij_atlas.out	512	512	0.030226
120	blocked_kij_atlas.out	1024	2	8.920727
121	blocked_kij_atlas.out	1024	4	1.823004
122	blocked_kij_atlas.out	1024	8	0.701317
123	blocked_kij_atlas.out	1024	16	0.509449
124	blocked_kij_atlas.out	1024	32	0.476489
125	blocked_kij_atlas.out	1024	64	0.269782
126	blocked_kij_atlas.out	1024	128	0.255873
127	blocked_kij_atlas.out	1024	256	0.245888
128	blocked_kij_atlas.out	1024	512	0.229296
129	blocked_kij_atlas.out	1024	1024	0.226316

2.4 BLAS dgemm routine

[6]:

	filename	matrix_size	timing
10	blas_routine.out	2	0.000023
11	blas_routine.out	4	0.000025
12	blas_routine.out	8	0.000027
13	blas_routine.out	16	0.000026
14	blas_routine.out	32	0.000039
15	blas_routine.out	64	0.000144
16	blas_routine.out	128	0.000635
17	blas_routine.out	256	0.004034
18	blas_routine.out	512	0.030143
19	blas_routine.out	1024	0.225817

3 Execution times of blocked algorithms in function of matrix size

We can plot the dependency of the execution on the matrix and block sizes in a 3d plot. I am using a log scale for the execution time since the differences in time are either really small or really big.

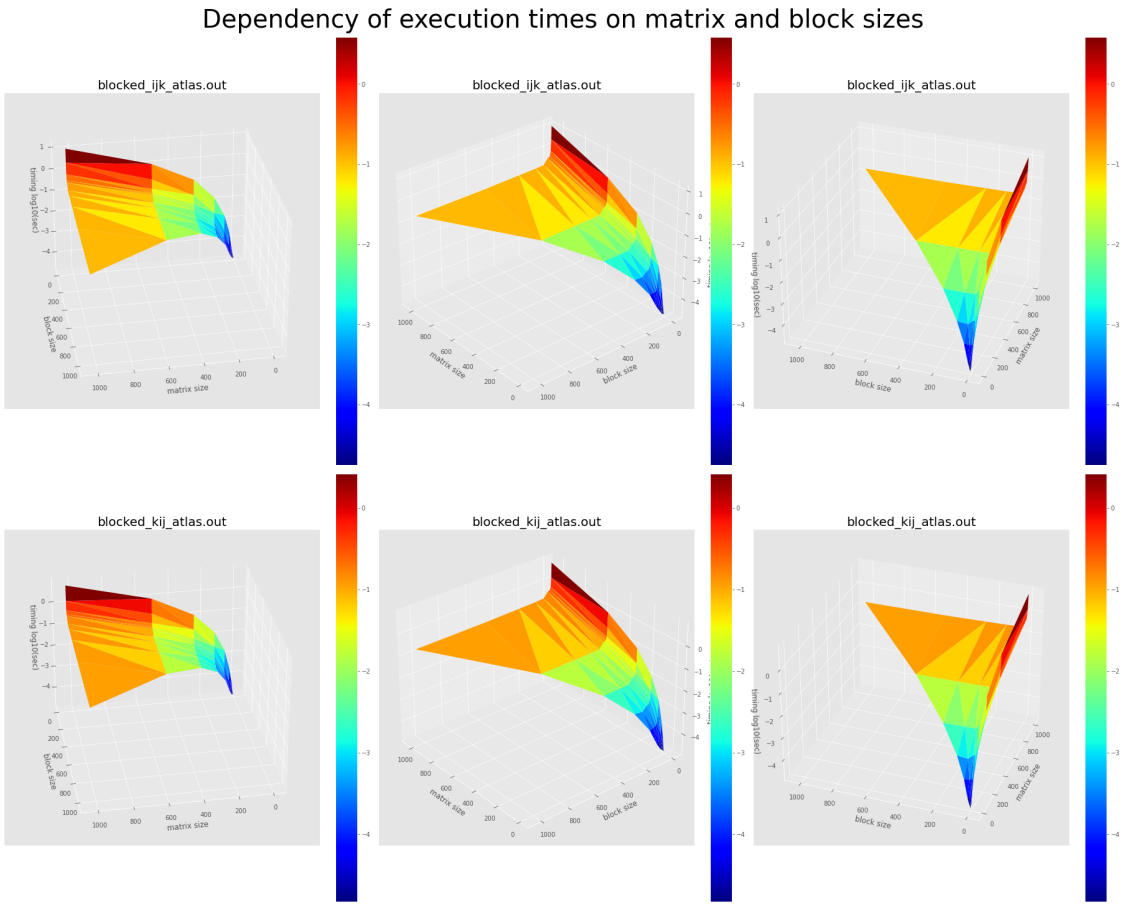

We can also plot in 2d for each matrix size:

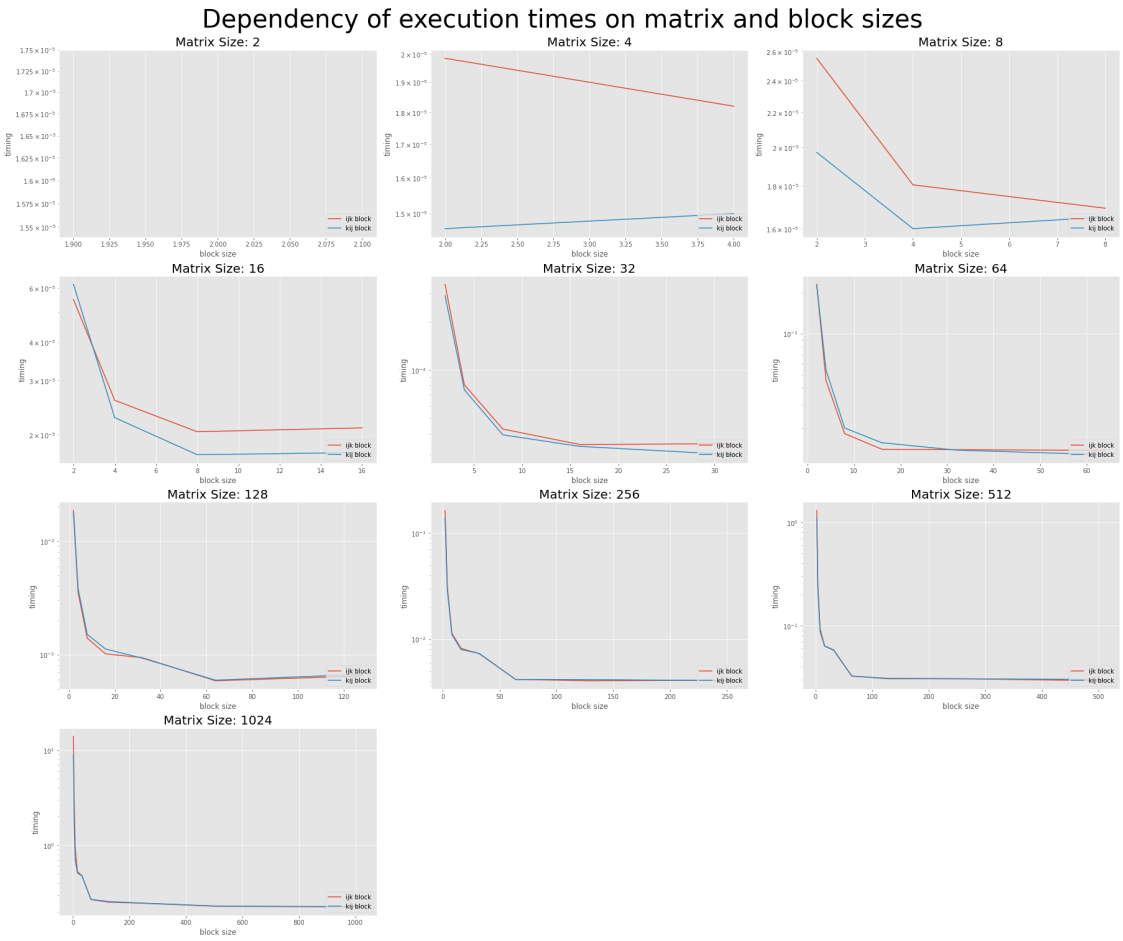

As we can see from these plots, the two blocked algorithm are similar in shape and in execution times. Blocking appears to not help with performance since there is no significant speedup when the block size is half the size of the matrix

3.1 Speedup of the blocked algorithms over the non-blocked one

Speedup is computed as following:

$$S(m,b) = \frac{T_{nonblocked}(m)}{T_{blocked}(m,b)}$$

where m is matrix size and b is block size

We can plot the speedup in 3d and 2d:

Speedup of the blocked algorithms over the non-blocked one

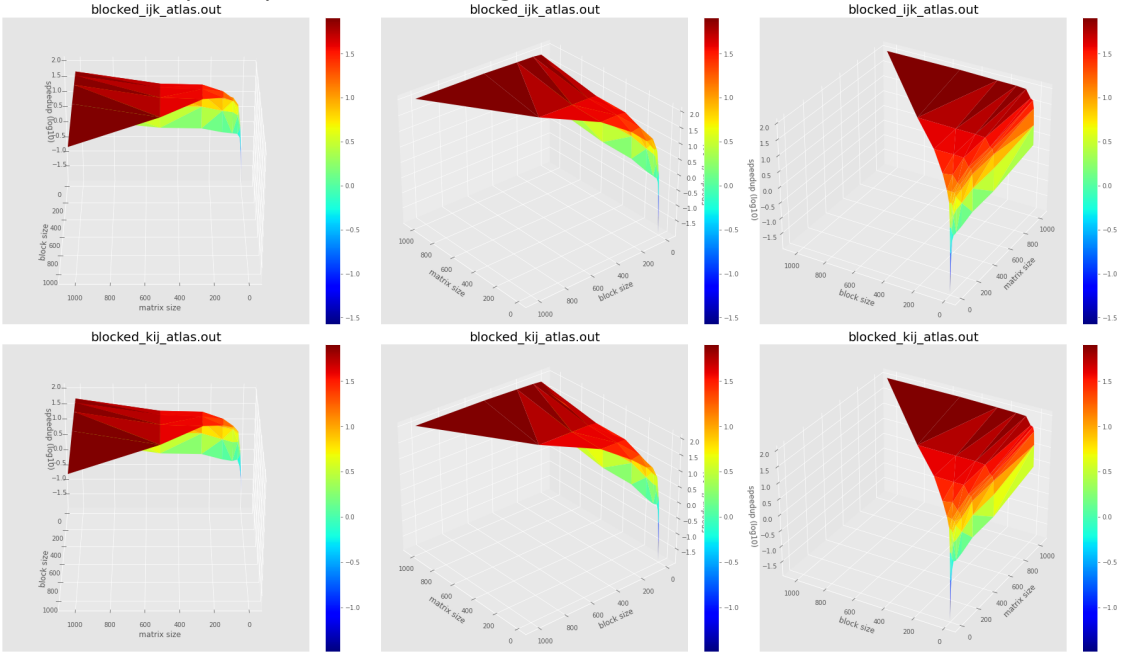

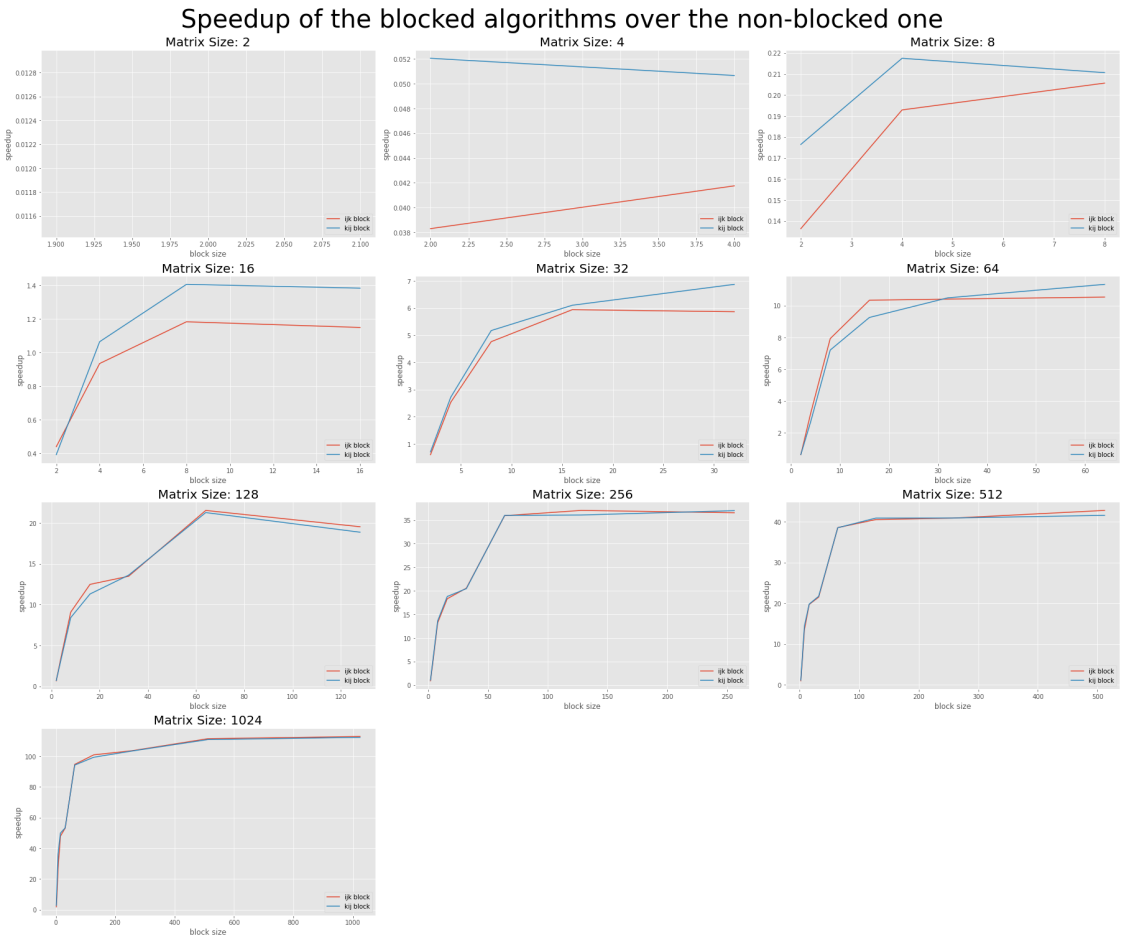

As we can see, the speedup can be up to a factor of 10^1 for large matrices ($m > 64$) and 10^2 for larger ones ($m > 1024$). For smaller matrix, we notice worse performance in comparison with the straightforward algorithm. There is no significant speedup between the ijk and kij algorithm any matrix size. However, I don't think we ca not conclude right now that blocked algorithms are faster than non-blocked ones since we are using ATLAS calls for the blocked algorithms which are already much faster than any naive implementation.

3.2 Comparison with BLAS dgemm

This time, we will compare our blocked algorithms with BLAS dgemm routine.

Speedup of the blocked algorithms over BLAS dgemm routine

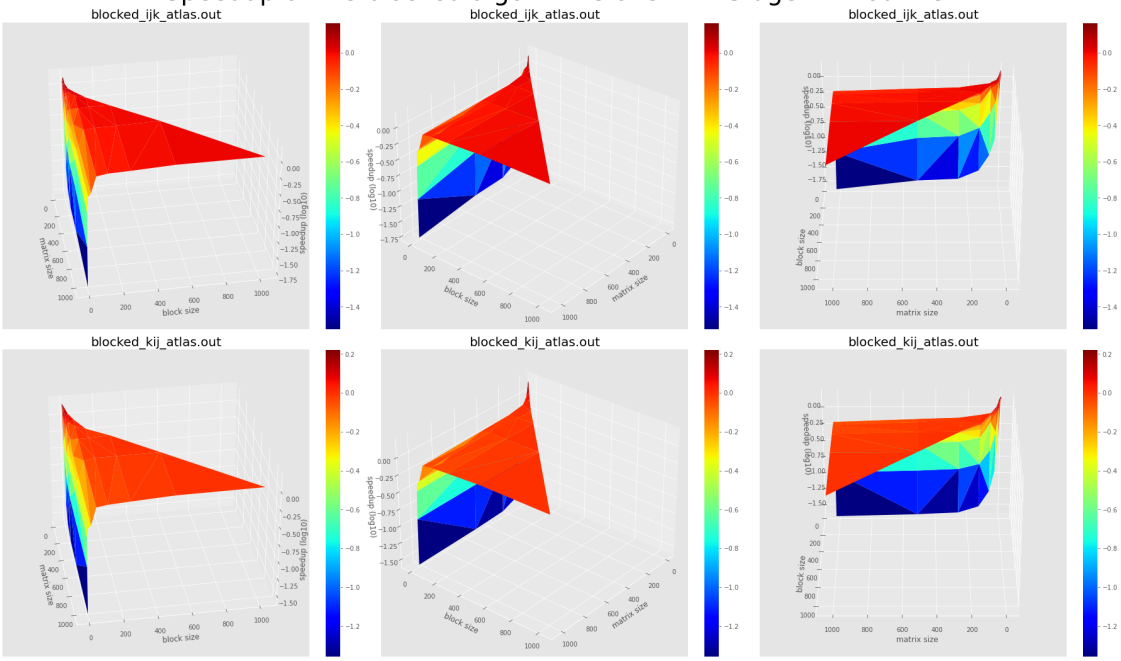

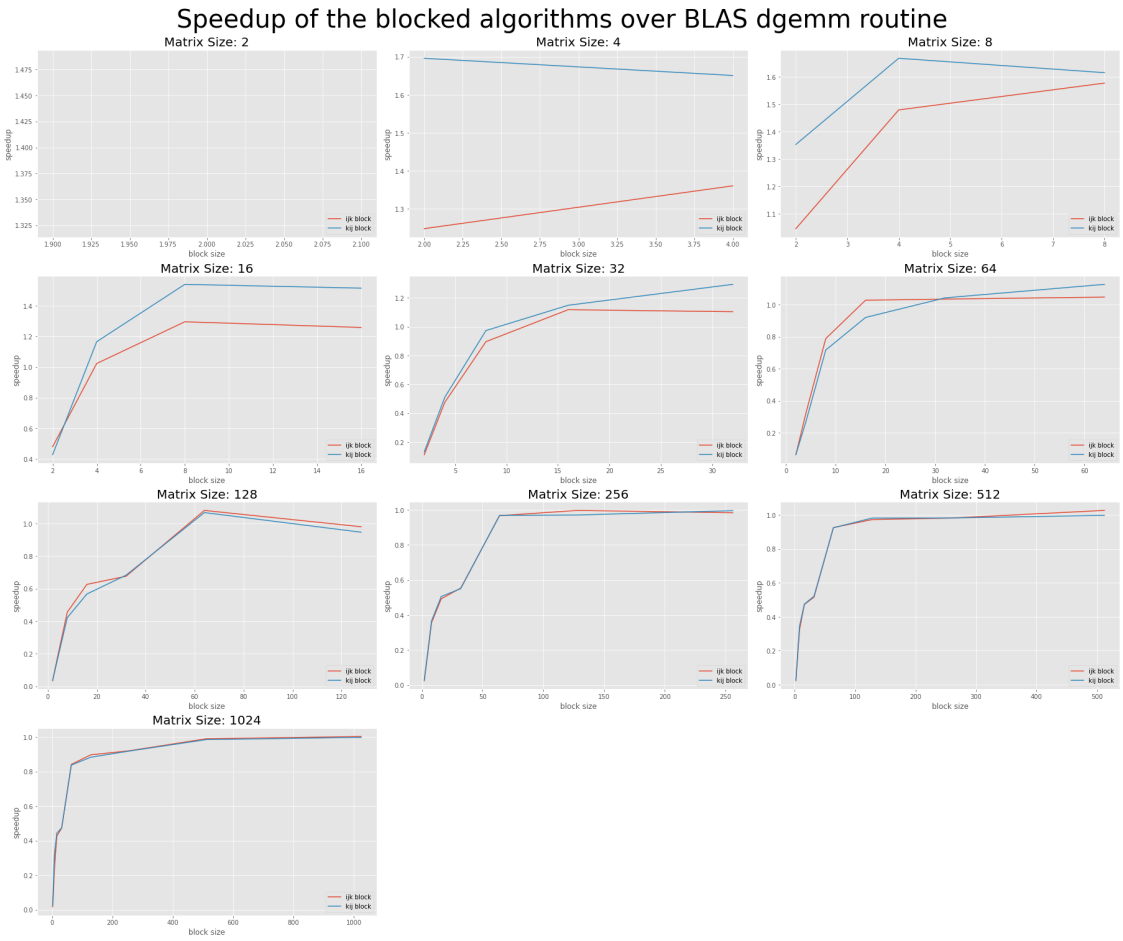

The speedup is generally under 1 for matrices with $m > 64$ and any block size. Blocking appears to lower the performance for matrix larger than 64x64 when using ATLAS calls. I expected blocking algorithms to give better performance since they reduce cache misses by reusing the same small blocks frequently. One reason I can think of why blocking in this case doesn't improve the performance is the `dgemm` routine's cost of overhead makes the multiple calls inefficient for small matrices. It means that for N_b number of blocks, we have $N_b.t_{overhead} + t_{procblocked} < t_{proc}$.